

# Matching Engine Specification

## Summary

A price-time priority limit order matching engine is responsible for determining when buy and sell orders transact (see Definitions below if unclear).

High frequency trading has pushed the limits of legacy matching engine implementations at the NYSE, CME, and other major exchanges. Proprietary trading groups can benefit from an internal matching engine to trim execution costs. Electronic market makers need to locally reconstruct the order book to determine when to provide liquidity. At the core of all these is the matching engine you will be writing.

Implement the interface provided in `engine.h` in a file `engine.c` for your entry. An example `engine.c` is provided. Rename or delete it before writing your own.

Language: C (gcc 4.4.5)

Starter Code: [match.zip](#)

Includes: API interface, autotester, makefile, simple example solution, type definitions, assumptions.

## Starter Code Descriptions

### engine.h

#### `init()`

Called once before the first orders are submitted. Initialize your data structures here. Set the next order id to 1.

#### `destroy()`

Called once after the last order is submitted. Free dynamically allocated memory here and cleanup. Calling `init()` after `destroy()` should completely reset the matching engine.

#### `limit()`

Called when a trader submits a new order. Check whether the order crosses with any existing orders. If so, call `execution`. If not, queue the order up with the others that have been submitted. Return the next order id and then increment it. The first order id returned is 1.

#### `cancel()`

Remove an order from the queue, identified by the order id returned from `limit()`.

#### `execution()`

Report that a trade occurred. Called twice whenever two orders cross, once with the buyer's name and side, and also with the seller's name and side. The two calls to `execution` may be in any order.

### types.h

t\_orderid

Order identification numbers which are returned by limit(). Used to uniquely identify the order for the sake of cancellation.

t\_price

Price of a share. In the range 0-65536, and interpreted as divided by 100. Eg the range is 000.00-655.36. Eg the price 123.45 = 12345; the price 23.45 = 2345; the price 23.4 = 2340

t\_size

Number of shares.

t\_side

Boolean representing whether a bid (0) or ask (1).

t\_order

Order submitted to the matching engine containing symbol, trader name, side, price, and size.

t\_execution

Execution report sent to trader informing them of the symbol, side, price, and size of the transaction.

## limits.h

MAX\_PRICE

Assume limit() will never be called with an order whose price field is greater than this number.

MIN\_PRICE

Assume limit() will never be called with an order whose price field is less than this number.

MAX\_LIVE\_ORDERS

Assume there will be no more than this number of uncrossed orders sitting on the book at any time.

STRINGLEN

Assume the trader names and order and execution symbols will be this long.

## Makefile

all

build the implementation in engine.c and the test script. Then call ./test to run the auto testing script.

clean

remove engine.o, executables, and emacs backups

## derived.h+c

Not required. These are provided to show you how one can derive interesting order types that are offered by many exchanges from the basic limit() and cancel() api.

## engine.c + double link list.c

An example implementation that is not optimized. It treats the bid and ask queues as doubly linked lists. This makes limit()  $O(n)$ , cancel()  $O(n)$  and in addition the constant factors, ignored by big- $O$ , are quite high.

### test.c

Script to automatically test your implementation's logic.

After implementing engine.c, call:

```
$ make all
```

```
$ ./test
```

Fix the errors before submission or your entry will be disqualified. Tests are incremental so they can be used to check implementation progress.

### score.c

Script to automatically test your implementation's speed. Shows how the scoring script works and what you should optimize for. This will be run on an isolated core (isolcpus=1).

After implementing engine.c, call:

```
$ make all
```

```
$ taskset 2 ./score
```

Try to minimize both the mean and the standard deviation of the latency.

### score\_feed.h

Simulated order and cancel data feed for use with score.c. This is an example of what the official scoring dataset will look like. Orders with price = 0 correspond to cancels with orderid = size.

## Scoring

Submit *only* your engine.c file by emailing it to maxdama at berkeley.edu with the subject line: QuantCup1:<Displayed Name>. I will verify <Displayed Name> with the email address you registered. Your registration information must be accurate to be scored. You may submit additional code but not submit modified versions of any of the provided code except for engine.c. For example one may submit engine.c and double\_link\_list.c since it does not require modification of Makefile.

Minimum Requirements to be Considered:

- 1) Implement all the matching engine logic
- 2) Compile without errors on Ubuntu 10.10 and run using dedicated core: isolcpus=1 and cpu affinity
- 3) Pass all the logic tests in the autotester script
- 4) Include comment explaining overall architecture and main optimizations

Score Formula = mean(latency) + sd(latency), as reported by ./score.

./score will use a different score\_feed.h than the one provided to discourage over-optimization.

The entry with the lowest score wins.

Disqualification:

You may be disqualified, without the judge providing justification nor entertaining appeals, for any perceived attempt to cheat or otherwise game the contest.

[Detailed information about judging platform](#)

## Suggestions

Data Structure:

[How to Build a Fast Limit Order Book \(copy\)](#) from WK Selph's HowToHFT Blog

Cache Optimization:

[Building a HPC Financial Exchange Video](#)

## Definitions

**Price-Time:** At any time there are multiple buyers waiting for someone to come sell shares at an acceptable price. Price-time priority means that the buyers will get filled in order of who offers the best price, and if the prices tie, then by who has been waiting the longest. The same prioritization applies to the sellers who submitted orders that did not fill immediately. Price-time is how the NYSE, NASDAQ, CME, and other major exchanges work.

**Matching Engine:** the computer program at the heart of an exchange which (1) stores all the outstanding limit orders (this data structure is called an "order book" or just "book"), (2) accepts new orders, (3) determines if new orders cross with any orders already on the book, (4) notifies both sides that their orders have been executed if it does cross, (5) queues orders in the book if they do not cross, (6) accepts order cancellation requests.

**Bid:** buy limit order

**Ask:** sell limit order

**Cross:** a bid and an offer having mutually acceptable prices. A cross results in a trade occurring and shares being transferred. For example a bid for 100 shares at \$100 on the book and a newly arriving ask for 50 shares at \$100 dollars cross at the price of \$100, resulting in the trade of 50 shares for \$100. The same trade would occur if the newly arriving ask has a price of \$99, or even less.

**Execution:** short for the execution report which is sent to a trader notifying them that a trade has occurred, how many shares have been transacted, and at what price. Executions are always sent in pairs, one to the buyer and one to the seller.

**Latency:** the time a computer process takes from start to finish. Can be broken down into matching engine, connection, network stack, data processing, and decision logic latency. In this challenge you will be optimizing matching engine latency only.

**Jitter:** roughly the standard deviation of latency. High jitter is bad because trades will nondeterministically be delayed.

## FAQ

**Q:** I assume that writing code in multiple files or some other language, compiling them, and sticking the source and a giant assembly dump in engine.c wouldn't be appreciated?

**A:** Correct. Code must be portably written in pure C.

**Q:** limits.h has the following: #define MAX\_LIVE\_ORDERS 65536. MAX\_TRADERS and MAX\_SYMBOLS must therefore be  $\leq 65536$ ... what would be a realistic upper limit for these?

**A:** MAX\_TRADERS = 50. MAX\_SYMBOLS = 1.